

Iterative Gradient based ~~Learning~~ Descent Pseudocode:

- 1) - Standardize data X of size $N \times D$
 - Target value matrix Y of size $N \times 1$
- 2) add bias colⁿ x so that $x = [1 \ x]$ & has size $N \times [D+1]$
- 3) initialize params θ_j for $j=1 \dots D+1$ to some random ~~data~~ ^{value}
 $\theta = \text{col}^n \text{ vector of size } (D+1) \times 1$
- 4) Until Convergence

$\forall x \in X$

- Compute gradient $\frac{dJ}{d\theta} = 2x^T(x\theta - y)$
- update $\theta \rightarrow \theta = \theta - \eta \frac{dJ}{d\theta}$

Termination Criteria
 We terminate when

- max iterations reached
- params change very little
- $\Delta \text{training error} = \text{very little}$

all this was iterative gradient descent \rightarrow considers 1 observation at a time
 problem \rightarrow gives local solⁿ & takes long time

* Solⁿ?

update params using array of gradients created by all observations
 batch gradient descent

$$\frac{d}{d\theta} \quad \frac{1}{N} \sum_{n=1}^N x^T (x\theta - y)$$

$$\theta = \theta - \frac{\eta}{N} x^T (x\theta - y)$$

let η absorb other scalar, 2

N \uparrow now we need all tho!!

Overfitting ??

* Solⁿ? mini batch.

- add regularization term
- penalty proportional to magnitude of θ ($\|\theta\| = \theta^T \theta$)
 $\lambda = \text{blending term}$
- stop using validation set: compute validation error on each step